

YOUR WEAK LLM IS SECRETLY A STRONG TEACHER FOR ALIGNMENT

Leitian Tao, Yixuan Li

Department of Computer Sciences, University of Wisconsin-Madison
 {leitiantao, sharonli}@cs.wisc.edu

ABSTRACT

The burgeoning capabilities of large language models (LLMs) have underscored the need for alignment to ensure these models act in accordance with human values and intentions. Existing alignment frameworks present constraints either in the form of expensive human effort or high computational costs. This paper explores a promising middle ground, where we employ a weak LLM that is significantly less resource-intensive than top-tier models, yet offers more automation than purely human feedback. We present a systematic study to evaluate and understand weak LLM’s ability to generate feedback for alignment. Our empirical findings demonstrate that weak LLMs can provide feedback that rivals or even exceeds that of fully human-annotated data. Our study indicates a minimized impact of model size on feedback efficacy, shedding light on a scalable and sustainable alignment strategy. To deepen our understanding of alignment under weak LLM feedback, we conduct a series of qualitative and quantitative analyses, offering novel insights into the quality discrepancies between human feedback vs. weak LLM feedback. Code is publicly available at https://github.com/deeplearning-wisc/weak_llm_teacher.

1 INTRODUCTION

As we observe the impressive capabilities of large language models (LLMs) across diverse applications (Brown et al., 2020; Achiam et al., 2023; Bubeck et al., 2023; Team et al., 2023; Anthropic, 2023), there emerges a critical need to ensure AI systems are helpful and harmless. AI alignment aims to harmonize AI behaviors with human intentions and values and ensure safe and desirable behavior. A key recipe to achieve alignment involves presenting pairs of responses and collecting binary feedback (e.g., preferred, less preferred) based on the comparative quality of these responses. The prevailing methods in alignment can be categorized based on the source of the feedback. For example, the popular framework Reinforcement Learning from Human Feedback (RLHF) (Christiano et al., 2017; Ziegler et al., 2019; Ouyang et al., 2022) relies on pure human judgments, which often involves considerable human labor and manual effort. On the other end of the spectrum, framework such as Reinforcement Learning from AI Feedback (RLAIF) (Bai et al., 2022a; Lee et al., 2023) harnesses feedback from high-capacity LLMs to annotate preference datasets, often incurring significant computational and financial costs, along with the need for heavy prompt engineering.

These contrasting frameworks highlight two extremes of a feedback spectrum, raising critical questions about the largely untapped middle ground that leverages the strengths of both while alleviating their respective drawbacks. As shown in Figure 1, this middle ground entails using significantly smaller LLMs that are less resource-intensive than top-tier models, yet offer more automation than purely human feedback. For instance, while a model like GPT-4 might have trillions of parameters, a weak LLM might only have hundreds of millions or even fewer. This smaller size means they inherently demand less computational power, thus reducing operational costs and enabling faster iteration cycles in development. This approach not only aligns with the practical demands of deploying scalable AI solutions but also addresses the need for sustainable development within AI research, making it a prudent choice for tasks where the ultra-high capabilities of state-of-the-art models do not necessarily translate to significant improvements in feedback quality.

Despite the appeal, *the research community still lacks a systematic evaluation and understanding of weak LLMs’ capability to generate feedback for alignment*. Motivated by this gap, we undertake a




Figure 1: A spectrum of feedback for aligning LLMs, ranging from labor-intensive human annotations (e.g., RLHF (Ouyang et al., 2022)) to highly automated, resource-intensive LLM feedback (e.g., RLAI (Bai et al., 2022a; Lee et al., 2023)). Our work explores the largely untapped middle ground, evaluating and understanding the use of weak LLM feedback for alignment.

comprehensive investigation of alignment via weak LLM feedback. Specifically, this paper makes three key contributions:

Contribution 1: A framework for evaluating alignment via weak LLM feedback. We formalize a learning and evaluation workflow utilizing feedback from weak LLM instead of traditional human annotations. Our framework operates on the combination of labeled and unlabeled preference datasets. Different from the existing RLHF framework, we leverage data comprising of unlabeled triplet (x, y_1, y_2) , where both y_1 and y_2 are responses corresponding to prompt x but the preference is unknown. In practice, the unlabeled triplets can be collected in large volumes without the need for human annotations. A weak LLM trained on labeled data will provide preference feedback on the large unlabeled data. Finally, we train a target LLM policy based on the weak LLM’s feedback. The framework introduces a novel perspective by connecting semi-supervised learning to alignment—a field that has yet to be thoroughly explored, especially in the context of using weak AI feedback.

Contribution 2: A comprehensive evaluation and novel empirical findings (Section 3). We employ the framework to evaluate the impact of weak LLM feedback on alignment across a variety of model scales and diverse model families. Intriguingly, *our results reveal that using a weak LLM, with size as small as 125M (Zhang et al., 2022), to provide preference label for alignment can match or even exceed the performance of using full human feedback* (see Figure 2 and Figure 3). Moreover, we systematically evaluate alignment performance when the feedback is provided by LLM of varying capacities: a weak supervisor (OPT-125M), a moderate supervisor (OPT-1.3B), a strong supervisor (Llama-3-8B), and a very strong supervisor (GPT-4). We found that the performance under weak, moderate, and strong supervisors is nearly comparable, suggesting that the supervisor model’s size has minimal impact on feedback effectiveness. Notably, the weak LLM, OPT-125M, outperforms the more advanced GPT-4, indicating that a task-specific weak LLM can provide more effective feedback than a larger, more powerful LLM that relies solely on prompt engineering.

Contribution 3: An in-depth analysis on the quality of weak LLM’s feedback (Section 4). To deepen our understanding of alignment under weak LLM feedback and reason our observations, we conduct a series of qualitative and quantitative analyses. A pivotal aspect of our study lies in the examination of quality discrepancies between human feedback vs. weak LLM feedback. Our key observations are threefold: (1) when weak LLM’s chosen response contradicts human’s choice, nearly half of these responses exhibit higher quality, suggesting the lack of reliability in human feedback and weak LLM can surpass human judgments; (2) weak LLM’s feedback can be qualitatively similar to human feedback in contexts with a clear gap, yet they exhibit uncertainty in judgments when the distinctions in response quality are subtle; and (3) advanced LLMs like GPT-4

also show increased feedback inconsistency in scenarios where response distinctions are minimal, highlighting the challenges faced by strong LLMs in providing feedback.

2 PRELIMINARIES ON LLM ALIGNMENT

We denote π_θ as a language model policy parameterized by θ , which takes in an input prompt x , and outputs a discrete probability distribution $\pi_\theta(\cdot|x)$ over the vocabulary space $\mathcal{V}$. $\pi_\theta(y|x)$ refers to the model’s probability of outputting response y given input prompt x . Alignment algorithms operate on comparative data, where pairs of responses are presented, and the model is trained to produce the preferred response given a query. Formally, we define the preference data below.

Definition 2.1 (Preference data) Consider two responses y_c, y_r for an input prompt x , we denote $y_c \succ y_r$ if y_c is preferred over y_r . We call y_c the chosen or preferred response and y_r the rejected response. Each triplet (x, y_c, y_r) is referred to as a preference. Furthermore, the empirical dataset $\mathcal{D} = \{(x_i, y_{c,i}, y_{r,i})\}_{i=1}^n$ consists of n such triplets sampled from a preference distribution.

Reinforcement Learning from Human Feedback. RLHF is a widely used paradigm for fine-tuning language models based on human preferences (Christiano et al., 2017; Ziegler et al., 2019; Ouyang et al., 2022; Bai et al., 2022a). The key stages in RLHF are reward modeling and reinforcement learning with the learned reward function.

Reward modeling learns a function mapping, which takes in the prompt x and response y and outputs a scalar value $r(x, y)$ signifying the reward. A preferred response should receive a higher reward, and vice versa. Based on the Bradley–Terry model (Bradley & Terry, 1952), the reward function is optimized over a dataset of human preferences, with the following objective:

$$\mathcal{L}_R = -\mathbb{E}_{(x, y_c, y_r) \in \mathcal{D}} [\log \sigma(r(x, y_c) - r(x, y_r))], \quad (1)$$

where σ denotes the sigmoid function. Using the learned reward function, the model is further fine-tuned with reinforcement learning to maximize the expected rewards, thus promoting exploration and adherence to learned preferences. The optimization objective is formulated as follows:

$$\max_{\pi_\theta} \mathbb{E}_{\hat{y} \sim \pi_\theta(\cdot|x)} [r(x, \hat{y})] - \beta \log \frac{\pi_\theta(\hat{y}|x)}{\pi_{\text{ref}}(\hat{y}|x)}, \quad (2)$$

where $\hat{y}$ represents the response generated by the current policy π_θ for the prompt x , π_{ref} indicates the reference policy or the initial policy before running RL optimization, and β serves as a hyperparameter to regulate the KL divergence.

Training with RLHF can be computationally expensive due to the use of multiple models. As an alternative, Rafailov et al. (2023) proposed to directly optimize for the policy best satisfying the preferences with a simple objective:

$$\mathcal{L}_{\text{DPO}}(\pi_\theta; \pi_{\text{ref}}; \mathcal{D}) = -\mathbb{E}_{(x, y_c, y_r) \in \mathcal{D}} \left[\log \sigma \left(\beta \left(\log \frac{\pi_\theta(y_c|x)}{\pi_{\text{ref}}(y_c|x)} - \log \frac{\pi_\theta(y_r|x)}{\pi_{\text{ref}}(y_r|x)} \right) \right) \right]. \quad (3)$$

Rafailov et al. (2023) showed that under mild assumptions, the optimal policy under the DPO objective (3) is equivalent to the optimal policy under the RLHF objective (2). This objective facilitates a more direct reflection of human preference judgments within the optimization framework.

3 HOW GOOD IS ALIGNMENT WITH WEAK LLM FEEDBACK?

In existing alignment approaches described above, models are often trained on fully supervised data where each preference label is hand-annotated by humans. To systematically evaluate the reliability of using feedback from a weak LLM instead of human feedback, we outline the training workflow in Section 3.1, experimental setup in Section 3.2, and present our main findings in Section 3.3.

3.1 ALIGNMENT VIA WEAK LLM FEEDBACK

Consider two empirical datasets, $\mathcal{D}_l$ and $\mathcal{D}_u$, representing a labeled preference dataset and an unlabeled dataset respectively. The labeled dataset $\mathcal{D}_l$ comprises of triplets (x, y_c, y_r) , where

human annotators have indicated a known preference $y_c \succ y_r$. $\mathcal{D}_u$, on the other hand, comprises of unlabeled triplet (x, y_1, y_2) , where both y_1 and y_2 are responses corresponding to x but the preference is unknown. A weak LLM, trained on $\mathcal{D}_l$, will provide preference feedback on the unlabeled data $\mathcal{D}_u$. This setup offers practical advantages since unlabeled triplets can be collected in large volumes without the need for human preference annotations. For example, one can generate multiple responses by querying varying LLMs with the same prompt. Below, we describe the alignment process using feedback from the weak LLM.

Preference feedback from the weak LLM. We first train a weak language policy π_w based on the labeled preference dataset $\mathcal{D}_l$. We use the subscript w to indicate “weak” in the remainder of the paper. Specifically, the weak LLM is optimized using the DPO loss (cf. Equation 3), under which the optimal policy is equivalent to that of RLHF:

$$\pi_w^* = \operatorname{argmin} \mathcal{L}_{\text{DPO}}(\pi_w; \pi_w^{\text{SFT}}; \mathcal{D}_l), \quad (4)$$

where π_w^* signifies the policy for the weak model, trained with DPO loss. π_w^{SFT} is the reference model or the initialization, which is an SFT model fine-tuned on the preferred question-answer pairs (x, y_c) in $\mathcal{D}_l$. Compared to directly using the untuned base model as a reference model, performing SFT enhances the model’s ability to generate desired responses to questions (see Appendix C).

We then generate the weak feedback for unlabeled data $\mathcal{D}_u$, leveraging the weak language model π_w^* . For each triplet $(x, y_1, y_2) \in \mathcal{D}_u$, we compute the reward $r_w(x, y_1)$ and $r_w(x, y_2)$ according to DPO’s implicit reward:

$$r_w(x, y) = \beta \log \frac{\pi_w(y|x)}{\pi_w^{\text{SFT}}(y|x)}. \quad (5)$$

We then assign the preference label $\hat{y}_c$ for the response with a higher predicted reward, and $\hat{y}_r$ for the response with a lower predicted reward. Mathematically:

$$\hat{y}_c = \begin{cases} y_1 & r_w(x, y_1) > r_w(x, y_2) \\ y_2 & r_w(x, y_1) \leq r_w(x, y_2) \end{cases}$$

We denote the resulting weakly labeled dataset as $\mathcal{D}_{\text{weak}} = \{(x, \hat{y}_c, \hat{y}_r)\}$, where $|\mathcal{D}_{\text{weak}}| = |\mathcal{D}_u|$.

Alignment with feedback from the weak LLM. Finally, we train an LLM policy π_θ based on weak LLM feedback $\mathcal{D}_{\text{weak}}$. The model is aligned using the following objective:

$$\pi_\theta^* = \operatorname{argmin} \mathcal{L}_{\text{DPO}}(\pi_\theta; \pi_\theta^{\text{SFT}}; \mathcal{D}_{\text{weak}}), \quad (6)$$

where π_θ^* is the resulting policy based on the weak LLM feedback. To obtain the reference model π_θ^{SFT} , we fine-tune the base model using question-answer pairs $(x, \hat{y}_c)$ in $\mathcal{D}_{\text{weak}}$. Under the training workflow, *a central unresolved question we address in this paper is how effectively LLMs can be aligned using feedback from weak LLMs as opposed to relying solely on human feedback*. The training workflow thus serves as the foundation to explore our core research question. To understand this, we conduct a systematic evaluation in the next section.

3.2 EXPERIMENTAL SETUP

Dataset. To evaluate the performance, we use the Anthropic HH-RLHF (Helpful and Harmless) dataset (Bai et al., 2022a), which is the most commonly used dataset for alignment. The dataset consists of 112,000 training samples and 12,500 test samples and is publicly available. Each sample includes a prompt and two responses, with one being preferred over the other. *The selected responses are annotated based on the opinions of crowd workers*, who assess which response is more helpful and harmless. We preprocess the dataset by filtering out samples with token lengths greater than 512, which yields 100,000 training samples and 11,000 test samples. We split the training data into two disjoint sets. The first subset is used as labeled data $\mathcal{D}_l$, and the remainder is used as the unlabeled data $\mathcal{D}_u$ (by disregarding the preference labels). We will vary the size of $\mathcal{D}_l$ in our ablation. We also evaluate on the Reddit TL;DR (TL;DR) summarization dataset from Stiennon et al. (2020), which consists of a Reddit post and several short summaries, judged for quality and informativeness by human evaluators.




Figure 2: (a) Alignment with feedback from a weak LLM (OPT-125M) can outperform human feedback. (b) Alignment performance on OPT-1.3B model under varying capability of supervisor. See Section 3.3 for details.

Models. For our main experiment, we utilize the OPT family model introduced by Zhang et al. (2022), which provides a spectrum of model sizes. This versatility enables us to evaluate the performance across different levels of model capability. Specifically, we employ OPT models of varying sizes (1.3B, 2.7B, 6.7B, and 13B) as our LLM policy models, trained with feedback from a weak LLM. Additionally, to validate the efficacy of our experiments, we consider more advanced open-sourced models, including Llama-2-7B (Touvron et al., 2023), Mistral-7B (Jiang et al., 2023a) and Gemma-7B (Team et al., 2024). For a comprehensive description of the hyper-parameters employed in our experiments, please refer to Appendix A.

Evaluation metrics. Given the test set $\mathcal{D}_{\text{test}}$, we evaluate the generation performance under two policies: π_{θ}^* (**with weak LLM feedback**), and π_h^* (**with human feedback**). π_h^* is trained with DPO loss on the same set of triplets in $\mathcal{D}_u$, except for using the original preference label provided by human annotators. This allows us to estimate the performance achieved with fully supervised data. For a fair comparison, π_{θ}^* and π_h^* always share the same model capacity and only differ in the source of feedback. We assess the generation performance using the following metrics:

- **Gold reward:** Previous studies (Gao et al., 2023; Coste et al., 2024; Xiong et al., 2023) commonly utilize gold reward as a metric for assessing the generation quality of language models. Gold reward is desirable due to the high cost associated with obtaining ground truth gold rewards from human annotators. We employ the output of a large auxiliary gold reward model, denoted as r_{gold} , to evaluate the quality of generated responses. For each test input prompt x from $\mathcal{D}_{\text{test}}$, we generate a response $\hat{y}$ according to a given language policy, and then compute its gold reward as $r_{\text{gold}}(x, \hat{y})$. A higher gold reward signifies that the model’s responses better align with desired preferences. Our results are consistent under alternative gold reward models; see Appendix C for details.
- **GPT-4 win-rate:** We employ GPT-4 as a proxy for human evaluation by instructing it to review two responses (from two different policies) to the same prompt. It then assigns a rating on a scale from 1 to 10. A higher win rate indicates that a policy on average produces more favorable answers.

3.3 MAIN RESULTS

Alignment with weak LLM feedback can outperform human feedback. In Figure 2(a), we evaluate the alignment performance using feedback from the weak LLM vs. human. Employing OPT-125M as our weak LLM to provide supervision (with the lowest capacity in the model family), we align student models of varying capacities (OPT-1.3B, OPT-2.7B, OPT-6.7B, and OPT-13B). Interestingly, the alignment performance using weak LLM feedback closely matches or even surpasses that of using human feedback. The similar alignment performance between π_{θ}^* and π_h^* has not been observed in previous studies. Moreover, we provide additional evaluation in Appendix C, and show that the observation holds for different weak LLMs such as GPT-Neo-125M (Gao et al., 2020) and Pythia-160M (Biderman et al., 2023). Given that the primary objective of alignment is to enhance the model’s ability to generate content that better resonates with human goals, our findings suggest that leveraging weak LLM feedback is a promising route without access to full human supervision. For an in-depth analysis of the unexpectedly high efficacy of weak LLM feedback, we defer to Section 4.




Figure 3: (a) Results on different model families. (b) GPT-4 evaluation for different models aligned with weak LLM feedback vs. human feedback.

Alignment performance under the varying capability of supervisor. Existing frameworks on learning from AI feedback, such as RLAI (Lee et al., 2023), employ more advanced LLMs (e.g., GPT-4) to annotate preference datasets for the alignment of student models. To understand the alignment performance under different supervisor models, we systematically consider a spectrum of supervisor LLMs of varying capabilities: (1) *weak supervisor* based on OPT-125M, (2) *moderate supervisor* based on OPT-1.3B, (3) *strong supervisor* based on Llama-3-8B, and (4) *very strong supervisor* based on GPT-4. For variants (1) through (3), we apply the same training workflow described in Section 3. For variant (4), we utilize GPT-4 as the LLM labeler, following the prompt suggested by Lee et al. (2023). Based on these four different sources of feedback, we compare the alignment performance on the same policy model OPT-1.3B. As shown in Figure 2 (b), the alignment performance under the weak, moderate, and strong supervisors is nearly comparable. This suggests that the **size of the supervisor model plays a less impactful role in the effectiveness of feedback**. Notably, we observe instances where the weak LLM with a small capacity, OPT-125M, can outperform GPT-4 in providing feedback on preferences. Although GPT-4 is a more advanced model, this finding indicates that a task-specific weak LLM can serve as a more effective supervisor than a larger, more powerful LLM that relies solely on prompt engineering.

3.4 ADDITIONAL ABLATIONS

Results on different model families. To further validate our findings, we extend our evaluation to additional model families, including Llama-2-7B (Touvron et al., 2023), Mistral-7B (Jiang et al., 2023a) and Gemma-7B (Team et al., 2024). We employ OPT-125M as the weak LLM to provide preference labels. Figure 3(a) displays the average gold rewards of the generated responses, based on policies π_{θ}^* (with weak LLM feedback) and π_h^* (with human feedback). The gold rewards achieved using weak LLM feedback consistently surpass those obtained from human feedback. This underscores the effectiveness of alignment through weak LLM feedback across various LLM families.

Results under GPT-4 evaluation. Beyond gold reward measurement, we assess the *win-rate* of the model aligned with weak LLM feedback (π_{θ}^*) against the human feedback (π_h^*), utilizing GPT-4 to judge the helpfulness and harmlessness within dialogue contexts. We randomly select 100 prompts from the test set of HH-RLHF and employed GPT-4 to determine whether responses generated under the policy π_{θ}^* were superior to those from π_h^* . Considering the diversity of model families and sizes, we evaluated the GPT-4 win-rate across five student model variants: OPT-1.3B, OPT-2.7B, OPT-6.7B, Llama-2-7B, and Mistral-7B. The weak LLM feedback all comes from the OPT-125m. As illustrated in Figure 3 (b), the win-rate for policy π_{θ}^* (with weak LLM feedback) approaches 50% for all cases, indicating that its performance competitively matches that of using human feedback π_h^* .

Results on different datasets and tasks. To examine the effectiveness of weak LLM feedback for alignment across broader NLP tasks, we turn our attention to the Reddit TL;DR dataset (Stiennon et al., 2020), which focuses on summarization—a stark contrast to dialogue tasks, which consists of a Reddit post and several short summaries, judged for quality and informativeness by human evaluators. Similar to the setup for HH-RLHF, we use OPT-125M as the weak LLM and assess the performance of weak LLM feedback v.s. human feedback for different student sizes: OPT-1.3B, OPT-2.7B,




Figure 4: (a) Results on different tasks (TL;DR). (b) Results under different data sizes.

OPT-6.7B, and OPT-13B. As depicted in Figure 4 (a), the average gold rewards for outputs sampled under the policies π_{θ}^* also outperform π_h^* across all student model sizes. This finding reinforces the efficacy of weak LLM feedback for alternative linguistic tasks.

Results under different data sizes. In this ablation, we systematically investigate the impact of dataset size. We vary the ratio between the labeled dataset size and the full dataset size: 1/16, 1/8, 1/4, 1/2, with the remaining data serving as the unlabeled subset. We train the weak LLM across different ratios with the same steps. We utilize OPT-125M as the weak LLM to align the policy model OPT-1.3B. Notably, as shown in Figure 4 (b), even under a small labeled dataset size with a ratio of 1/16, the performance using weak LLM feedback is favorably close to that of using human feedback.

Section 3 key takeaways

1. Using a weak LLM, with size as small as 125M, to provide preference feedback for alignment can match or even exceed the performance of using pure human feedback.
2. Alignment performance can be similar using feedback from supervisor LLMs of varying capabilities—from weak (125M), moderate (1.3B), strong (8B) to very strong (GPT-4). **This suggests that the size of the supervisor LLM plays a less impactful role** in providing feedback for alignment.
3. Our finding consistently holds across different model families, evaluation metrics, and tasks.

4 IN-DEPTH UNDERSTANDING OF WEAK LLM’S FEEDBACK

When weak LLM feedback matches/mismatches human feedback. In this ablation, we investigate the impact of samples either matching or mismatching human feedback. Adopting OPT-125M as the weak LLM to provide feedback for the OPT-1.3B model alignment, we design three controlled settings for this ablation:

- $\mathcal{D}_{\text{weak}}$: Use all samples, $\{(x, \hat{y}_c, \hat{y}_r)\}$, for training. This is the same as our main setting.
- $\mathcal{D}_{\text{weak}}^{\text{match}}$: Use samples with weak labels *matching* the human feedback, i.e., $(r_w(x, y_c) > r_w(x, y_r))$.
- $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$: Use samples with labels *mismatching* the human feedback, i.e., $(r_w(x, y_c) < r_w(x, y_r))$.

As shown in Table 1, training with $\mathcal{D}_{\text{weak}}^{\text{match}}$ (samples with matching labels between weak LLM feedback and human feedback) yields gold rewards comparable to those obtained from using the full set $\mathcal{D}_{\text{weak}}$. Additionally, when examining the alignment performance using $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$, the gold reward significantly improves relative to the pre-trained strong model without any fine-tuning by 1.79. This indicates that $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$ effectively improves alignment, even when the weak LLM provides feedback that are entirely contrary to the human preference. We further confirm that the noise of preference labels in $\mathcal{D}_{\text{weak}}$ appears to behave differently from random noise (see Appendix C). The phenomenon is intriguing and prompts our next analysis.

Table 1: Effect of purifying weak LLM feedback. Numbers in brackets denote improvement relative to the pre-trained student model without any fine-tuning.

	% matching human feedback	Gold reward of aligned model π_{θ}^*
$\mathcal{D}_{\text{weak}}$	60.6%	4.84 (+2.61)
$\mathcal{D}_{\text{weak}}^{\text{match}}$	100%	4.78 (+2.56)
$\mathcal{D}_{\text{weak}}^{\text{mismatch}}$	0%	4.01 (+1.79)

Table 3: Qualitative example of the chosen and rejected responses in $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$.

Prompt
Human: What kind of vegetables offer the most health benefits?
Assistant:
Chosen response by weak LLM (rejected by human)
I'm not a doctor, but my guess is that kale, broccoli, cauliflower, and cabbage are good sources of vitamins and fiber. (gold reward: 6.82)
Rejected response by weak LLM (chosen by human)
Most health benefits, or more specifically, most nutrients? (gold reward: 3.25)

Dive into the quality of weak LLM vs. human feedback. In Table 3, we provide a qualitative example from $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$ where the chosen response by weak LLM mismatches that of the human feedback. More qualitative examples are in Appendix B. Despite contradicting the human preference, the chosen response by the weak LLM has higher quality. We further measure the gold reward as a proxy for response quality, which is indeed higher for the response chosen by the weak LLM, i.e., $r_{\text{gold}}(x, \hat{y}_c) > r_{\text{gold}}(x, \hat{y}_r)$. Across all the samples in $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$, we notice that 44.3% samples display a higher gold reward for the response chosen by the weak LLM. *This indicates that the preferences from weak LLM in $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$ are not entirely erroneous, but better than human feedback in nearly half of the cases¹.* Considering the feedback from weak LLM and human shares $\mathcal{D}_{\text{weak}}^{\text{match}}$ portion and only differs in $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$, the close performance achieved between the two models can be attributed to the fact that human feedback is similarly imperfect.

Building on the observation that weak LLM can sometimes surpass human judgment, we delve into the characteristics of datasets produced under such supervision. Table 2 shows that the gold rewards for chosen responses in $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$ are on average comparable to those for rejected responses, while the gold rewards for chosen responses in $\mathcal{D}_{\text{weak}}^{\text{match}}$ are significantly higher than the rejected responses. *This suggests that the weak LLM primarily errs between subtly different choices, while it remains reliable in distinguishing between options with clear quality distinctions.* Overall, the discrepancy (Δ) within $\mathcal{D}_{\text{weak}}$ is similar to that in $\mathcal{D}_{\text{human}}$, indicating that the quality of weak labels is not severely compromised.

More qualitative analysis. To further investigate the quality gap in responses from the dataset where weak LLM feedback and human feedback diverge ($\mathcal{D}_{\text{weak}}^{\text{mismatch}}$), we conducted a qualitative assessment using GPT-4. Our goal is to determine whether GPT-4’s preferences for chosen and rejected responses in $\mathcal{D}_{\text{weak}}^{\text{match}}$ and $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$ remain consistent. We hypothesize that consistent preferences from GPT-4 across multiple evaluations would indicate a significant quality difference between the responses, while fluctuating preferences would suggest that the responses are less distinguishable in quality. To quantify this, we prompt GPT-4 ten times consecutively to identify which answer was more helpful and less harmful. We randomly sample 100 examples from $\mathcal{D}_{\text{weak}}^{\text{match}}$ and $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$, and measure preference consistency by calculating the fraction of majority votes for each question, which ranges from 0.5 to 1. For instance, Table 5 presents a qualitative example from $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$, where the weak LLM and humans have different preferences. We observe that GPT-4 indeed exhibits mixed preferences across the ten evaluations, indicating a low preference consistency (0.5) and thus the indistinguishable nature between the two responses. Indeed, as illustrated in Table 4, the average consistency for

Table 2: Summary statistics of average gold reward on HH-RLHF dataset. “Chosen” and “Rejected” indicate the human preference for original dataset $\mathcal{D}_{\text{human}}$ and the weak LLM preference for the $\mathcal{D}_{\text{weak}}$, $\mathcal{D}_{\text{weak}}^{\text{match}}$ and $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$.

	Chosen	Rejected	Δ
$\mathcal{D}_{\text{human}}$	5.77	4.23	1.52
$\mathcal{D}_{\text{weak}}$	5.63	4.39	1.23
$\mathcal{D}_{\text{weak}}^{\text{match}}$	5.93	3.66	2.27
$\mathcal{D}_{\text{weak}}^{\text{mismatch}}$	5.17	5.53	-0.36

Table 4: The GPT-4 preference consistency.

Dataset	consistency
$\mathcal{D}_{\text{weak}}^{\text{match}}$	0.84
$\mathcal{D}_{\text{weak}}^{\text{mismatch}}$	0.66

¹This is further evidenced by our GPT-4 evaluation, where the response chosen by weak LLM achieved a 46% win rate over the rejected response.

Table 5: A qualitative example in $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$ with low preference consistency in GPT-4 eval.

Prompt										
<i>Human:</i> I'm visiting the tropics and I'd like to go snorkeling in coral reefs to see fish, but I've never done it before. What equipment do I need?										
<i>Assistant:</i> Snorkeling is very common among tourists. You can probably just show up at the beach and rent a snorkeling mask, snorkel, fins, and a life jacket for a very small fee. I'd recommend you try snorkeling on calm, flat water as it's more relaxing.										
<i>Human:</i> Don't life jackets prevent people from going underwater? Why would I wear one while snorkeling underwater?										
<i>Assistant:</i>										
Chosen response by weak supervisor (rejected by human)										
A life jacket is an inflatable vest that helps you float on the water. It's not supposed to prevent you from going underwater.										
Rejected response by weak supervisor (chosen by human)										
Right, I was just going to say that you might feel more comfortable wearing a life jacket until you get comfortable with the water. I've found that even the most advanced snorkelers prefer to wear life jackets in the beginning.										
GPT-4 preferred answer with 10 consecutive trials (consistency: 0.5)										
First	Second	First	First	Second	Second	First	Second	Second	Second	First

samples in $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$ is lower than that of $\mathcal{D}_{\text{weak}}^{\text{match}}$. This suggests that advanced LLMs such as GPT-4 experience significantly higher inconsistency when the quality distinctions between two responses are minimal, highlighting the challenges faced by both weak and ultra-large LLMs in providing feedback.

Section 4 key takeaways

1. When weak LLM's chosen response contradicts human feedback, nearly half of these responses exhibit higher gold rewards, suggesting that weak LLM can sometimes surpass human judgments.
2. Weak LLM primarily errs between choices that are subtly different, while it remains reliable in distinguishing between options with clear quality distinctions.
3. Feedback from advanced LLM (e.g., GPT-4) can exhibit high inconsistency in feedback when the distinctions between responses are subtle.

5 DISCUSSION

Differences w.r.t. Burns et al. (2024). Previous work by Burns et al. (2024) explored weak-to-strong generalization, wherein a less capable model's weak supervision signals are used to guide a stronger, larger model. However, *their study was limited to much simpler learning tasks like reward modeling and binary classification*, where the outputs are either scalar or categorical labels. They did not investigate the more challenging task of language generation, which is closely tied to alignment and involves a significantly more complex output space. Hence, the potential of weak LLM feedback for alignment remains unexplored in Burns et al. (2024), which is the novel focus of our study.

Contrary to Burns et al. (2024), which employed discriminative metrics such as accuracy, our study evaluates the generative performance essential for AI alignment. Although Burns et al. (2024) reported a noticeable performance gap between models trained with weak LLM feedback versus those with human feedback, our findings challenge these conclusions. **Considering that human preferences are noisy and unreliable** (Yeh et al., 2024), accuracy on the preference test set is a questionable metric. Our works show contrary conclusions when directly measuring the generation performance of the aligned model, and demonstrate that using feedback from weak LLM can not only match but potentially exceed human feedback in alignment tasks. Our work suggests a promising direction for future research in leveraging weak AI signals for alignment.

6 RELATED WORK

Large language model alignment. The primary objective of model alignment is to steer language models toward human-desired outputs. Numerous studies have leveraged human feedback to refine language models by human preferences (Christiano et al., 2017; Ziegler et al., 2019; Stiennon et al., 2020; Lee et al., 2021; Ouyang et al., 2022; Nakano et al., 2022; Glaese et al., 2022; Snell et al., 2023; Yuan et al., 2023; Song et al., 2024; Dong et al., 2023; Bai et al., 2022b; Lee et al., 2024; Munos et al., 2024; Hejna et al., 2024; Dai et al., 2024; Khanov et al., 2024). However, human preference data can be costly to collect and suffer from unreliability issues (Yeh et al., 2024), prompting researchers to explore AI-generated feedback for alignment purposes (Bai et al., 2022a; Lee et al., 2023; Ding et al.,

2023; Gilardi et al., 2023; Guo et al., 2024). Given the computational inefficiency of RLHF, recent shifts in focus favor closed-form losses that directly utilize offline preferences, like Direct Preference Optimization (Rafailov et al., 2023) and related methodologies (Gheshlaghi Azar et al., 2024; Pal et al., 2024; Liu et al., 2024b; Xiong et al., 2023; Tang et al., 2024; Yu et al., 2024; Ethayarajh et al., 2024; Zeng et al., 2024; Calandriello et al., 2024; Muldrew et al., 2024; Ray Chowdhury et al., 2024; Liu et al., 2024a; Gao et al., 2024; Yang et al., 2024; Chakraborty et al., 2024; Zhao et al., 2023). Beyond algorithmic approaches, researchers recently have provided theoretical understandings of the learning dynamics (Im & Li, 2024b) and generalization guarantee (Im & Li, 2024a) of DPO. Nonetheless, both RLHF and offline preference-based methods presuppose access to high-quality human, which limits their scalability. Our research endeavor takes initial steps toward understanding the effectiveness of aligning language models under weak LLM feedback. We present a systematic evaluation and in-depth analysis that sheds light on the quality and feasibility of leveraging weak LLM feedback for alignment, and reveal novel empirical findings to the research community.

Large language model as a judge. The use of LLM-as-a-Judge prompting to evaluate language models has become a common practice (Dubois et al., 2023; Li et al., 2023; Bai et al., 2023; Saha et al., 2023). This approach is also used in collecting preference datasets for alignment. For example, Bai et al. (2022a) employs an LLM to assess responses, refine them, and then use the resulting data to train a reward model known as “RL from AI Feedback” (RLAIF). Similarly, Lee et al. (2023) demonstrated that employing a strong LLM for LLM-as-a-Judge prompting to create preference datasets yields performance nearly on par with traditional RLHF. Yuan et al. (2024) explored the concept of LLM-as-a-Judge prompting to enable models to generate their own rewards during training. While many studies (Jiang et al., 2023b; Li et al., 2024b; Kim et al., 2023; Chen et al., 2024; Ye et al., 2024) focus on leveraging advanced LLMs as judges for preference dataset collection, we investigate the effectiveness of feedback from weak LLMs with significantly smaller capacity. We reveal new insight that a task-specific weak LLM can provide more effective preference feedback than an ultra-large LLM (*e.g.*, GPT-4) that relies solely on prompt engineering.

Weak-to-strong generalization. Weak-to-strong generalization refers to the scenario where a weaker teacher model supervises a stronger student model. Unlike traditional teacher-student frameworks such as semi-supervised learning (Laine & Aila, 2017; Tarvainen & Valpola, 2017), domain adaptation (French et al., 2018; Chen et al., 2019), and knowledge distillation (Hinton et al., 2015; Beyer et al., 2022)—which generally involve a stronger teacher guiding a weak student—other studies have explored employing comparably or even lesser-capable teachers to guide student model (Freund & Schapire, 1997; Furlanello et al., 2018; Xie et al., 2020; Higuchi et al., 2020; Burns et al., 2024; Green Larsen & Ritzert, 2022; Lang et al., 2024; Charikar et al., 2024). Recognizing the broad generalization capacities of large-scale pre-trained models, Burns et al. (2024) introduced the concept of weak-to-strong generalization for LLM, focusing on simpler discriminative tasks such as reward modeling rather than generative tasks such as alignment. A detailed discussion on the differences with Burns et al. (2024) is in Section 5. Different from previous works that explore the effectiveness of weak LLM under the standard supervised fine-tuning (Liu & Alahi, 2024; Li et al., 2024a; Ji et al., 2024; Hase et al., 2024; Sun et al., 2024; Bansal et al., 2024), we aim to delve deeper into the role of weak LLM could play for alignment. These tasks pose unique challenges due to inherently diverse output and the complexity of performance evaluation but hold promise for advancing superalignment in generative models (Puthumanailam et al., 2024). Instead of adopting model interpolation for alignment without any additional training (Zheng et al., 2024; Zhou et al., 2024), we explore the three-stage framework for alignment with feedback from weak LLM.

7 CONCLUSION

In this paper, we systematically explore and evaluate the effectiveness of leveraging weak LLM feedback for alignment. Our findings reveal that using a weak LLM, with a size as small as 125M, to provide preference feedback for alignment can match or even exceed the performance of using pure human feedback. Through our in-depth analyses, we shed light on the intricacies of alignment under weak LLM feedback and offer valuable insights into the quality discrepancies between human feedback vs. weak LLM feedback. This study not only highlights the feasibility of achieving high-quality alignment with less precise preference data but also suggests directions for future research to further refine alignment methods, ensuring AI systems remain beneficial and aligned with human intentions. We discuss the limitations and implications for future work in Section D.

ACKNOWLEDGEMENT

Research is supported in part by the AFOSR Young Investigator Program under award number FA9550-23-1-0184, National Science Foundation (NSF) Award No. IIS-2237037 & IIS-2331669, and the UL Research Institutes through the Center for Advancing Safety of Machine Intelligence. The authors would like to thank Shawm Im and Wendi Li for the valuable feedback on the work.

REFERENCES

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Anthropic. Claude: Conversational ai assistant, 2023.
- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022a.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022b.
- Yushi Bai, Jiahao Ying, Yixin Cao, Xin Lv, Yuze He, Xiaozhi Wang, Jifan Yu, Kaisheng Zeng, Yijia Xiao, Haozhe Lyu, Jiayin Zhang, Juanzi Li, and Lei Hou. Benchmarking foundation models with language-model-as-an-examiner. In *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2023.
- Hritik Bansal, Arian Hosseini, Rishabh Agarwal, Vinh Q Tran, and Mehran Kazemi. Smaller, weaker, yet better: Training llm reasoners via compute-optimal sampling. *arXiv preprint arXiv:2408.16737*, 2024.
- Lucas Beyer, Xiaohua Zhai, Amélie Royer, Larisa Markeeva, Rohan Anil, and Alexander Kolesnikov. Knowledge distillation: A good teacher is patient and consistent. In *Computer Vision and Pattern Recognition*, pp. 10925–10934, 2022.
- Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, et al. Pythia: A suite for analyzing large language models across training and scaling. In *International Conference on Machine Learning*, pp. 2397–2430, 2023.
- Ralph Allan Bradley and Milton E Terry. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33:1877–1901, 2020.
- Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, et al. Sparks of artificial general intelligence: Early experiments with gpt-4. *arXiv preprint arXiv:2303.12712*, 2023.
- Collin Burns, Pavel Izmailov, Jan Hendrik Kirchner, Bowen Baker, Leo Gao, Leopold Aschenbrenner, Yining Chen, Adrien Ecoffet, Manas Joglekar, Jan Leike, Ilya Sutskever, and Jeffrey Wu. Weak-to-strong generalization: Eliciting strong capabilities with weak supervision. In *International Conference on Machine Learning*, pp. 4971–5012, 2024.
- Daniele Calandriello, Zhaohan Daniel Guo, Remi Munos, et al. Human alignment of large language models through online preference optimisation. In *Proceedings of the 41st International Conference on Machine Learning*, pp. 5409–5435, 2024.
- Souradip Chakraborty et al. Maxmin-rlhf: Alignment with diverse human preferences. In *Forty-first International Conference on Machine Learning*, 2024.

- Moses Charikar, Chirag Pabbaraju, and Kirankumar Shiragur. Quantifying the gain in weak-to-strong generalization. *Advances in Neural Information Processing Systems*, 2024.
- Chaoqi Chen, Weiping Xie, Wenbing Huang, Yu Rong, Xinghao Ding, Yue Huang, Tingyang Xu, and Junzhou Huang. Progressive feature alignment for unsupervised domain adaptation. In *Computer Vision and Pattern Recognition*, pp. 627–636, 2019.
- Lichang Chen, Shiyang Li, Jun Yan, Hai Wang, Kalpa Gunaratna, Vikas Yadav, Zheng Tang, Vijay Srinivasan, Tianyi Zhou, Heng Huang, et al. AlpaGasus: Training a better alpaca with fewer data. In *The Twelfth International Conference on Learning Representations*, 2024.
- Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. *Advances in Neural Information Processing Systems*, 30, 2017.
- Thomas Coste, Usman Anwar, Robert Kirk, and David Krueger. Reward model ensembles help mitigate overoptimization. *International Conference on Learning Representations*, 2024.
- Josef Dai et al. Safe rlhf: Safe reinforcement learning from human feedback. In *The Twelfth International Conference on Learning Representations*, 2024.
- Bosheng Ding, Chengwei Qin, Linlin Liu, Yew Ken Chia, Boyang Li, Shafiq Joty, and Lidong Bing. Is GPT-3 a good data annotator? In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (eds.), *Proceedings of the Annual Meeting of the Association for Computational Linguistics*, 2023.
- Hanze Dong, Wei Xiong, Deepanshu Goyal, Yihan Zhang, et al. Raft: Reward ranked finetuning for generative foundation model alignment. *Transactions on Machine Learning Research*, 2023.
- Yann Dubois, Xuechen Li, Rohan Taori, Tianyi Zhang, Ishaan Gulrajani, Jimmy Ba, Carlos Guestrin, Percy Liang, and Tatsunori B Hashimoto. AlpacaFarm: A simulation framework for methods that learn from human feedback. *arXiv preprint arXiv:2305.14387*, 2023.
- Kawin Ethayarajh, Winnie Xu, Niklas Muennighoff, et al. Model alignment as prospect theoretic optimization. In *Proceedings of the 41st International Conference on Machine Learning*, pp. 12634–12651, 2024.
- Geoffrey French, Michal Mackiewicz, and Mark Fisher. Self-ensembling for visual domain adaptation. In *International Conference on Learning Representations*, 2018.
- Yoav Freund and Robert E Schapire. A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of computer and system sciences*, 55(1):119–139, 1997.
- Tommaso Furlanello, Zachary Lipton, Michael Tschannen, Laurent Itti, and Anima Anandkumar. Born again neural networks. In *International Conference on Machine Learning*, pp. 1607–1616, 2018.
- Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, Horace He, Anish Thite, Noa Nabeshima, et al. The pile: An 800gb dataset of diverse text for language modeling. *arXiv preprint arXiv:2101.00027*, 2020.
- Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. In *International Conference on Machine Learning*, pp. 10835–10866, 2023.
- Songyang Gao et al. Linear alignment: A closed-form solution for aligning human preferences without tuning and feedback. In *Forty-first International Conference on Machine Learning*, 2024.
- Mohammad Gheshlaghi Azar et al. A general theoretical paradigm to understand learning from human preferences. *AISTATS*, 2024.
- Fabrizio Gilardi, Meysam Alizadeh, and Maël Kubli. Chatgpt outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences*, 2023.
- Amelia Glaese, Nat McAleese, Maja Trebacz, John Aslanides, Vlad Firoiu, Timo Ewalds, Maribeth Rauh, Laura Weidinger, Martin Chadwick, Phoebe Thacker, et al. Improving alignment of dialogue agents via targeted human judgements. *arXiv preprint arXiv:2209.14375*, 2022.

- Kasper Green Larsen and Martin Ritzert. Optimal weak to strong learning. *Advances in Neural Information Processing Systems*, 35:32830–32841, 2022.
- Shangmin Guo, Biao Zhang, Tianlin Liu, Tianqi Liu, Misha Khalman, Felipe Llinares, Alexandre Rame, Thomas Mesnard, Yao Zhao, Bilal Piot, et al. Direct language model alignment from online ai feedback. *arXiv preprint arXiv:2402.04792*, 2024.
- Peter Hase, Mohit Bansal, Peter Clark, and Sarah Wiegrefe. The unreasonable effectiveness of easy training data for hard tasks. *arXiv preprint arXiv:2401.06751*, 2024.
- Joey Hejna et al. Contrastive preference learning: Learning from human feedback without reinforcement learning. In *The Twelfth International Conference on Learning Representations*, 2024.
- Yosuke Higuchi, Shinji Watanabe, Nanxin Chen, Tetsuji Ogawa, and Tetsunori Kobayashi. Mask ctc: Non-autoregressive end-to-end asr with ctc and mask predict. *Annual Conference of the International Speech Communication Association*, 2020.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Shawn Im and Yixuan Li. On the generalization of preference learning with dpo. *arXiv preprint 2408.03459*, 2024a.
- Shawn Im and Yixuan Li. Understanding the learning dynamics of alignment with human feedback. In *International Conference on Machine Learning*, 2024b.
- Jiaming Ji, Boyuan Chen, Hantao Lou, Donghai Hong, Borong Zhang, Xuehai Pan, Juntao Dai, and Yaodong Yang. Aligner: Achieving efficient alignment through weak-to-strong correction. *arXiv preprint arXiv:2402.02416*, 2024.
- Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. Mistral 7b. *arXiv preprint arXiv:2310.06825*, 2023a.
- Dongfu Jiang, Xiang Ren, and Bill Yuchen Lin. Llm-blender: Ensembling large language models with pairwise ranking and generative fusion. *arXiv preprint arXiv:2306.02561*, 2023b.
- Maxim Khanov, Jirayu Burapachee, and Yixuan Li. Args: Alignment as reward-guided search. In *Proceedings of the International Conference on Learning Representations*, 2024.
- Seungone Kim, Jamin Shin, Yejin Cho, Joel Jang, Shayne Longpre, Hwaran Lee, Sangdoo Yun, Seongjin Shin, Sungdong Kim, James Thorne, et al. Prometheus: Inducing fine-grained evaluation capability in language models. *arXiv preprint arXiv:2310.08491*, 2023.
- Samuli Laine and Timo Aila. Temporal ensembling for semi-supervised learning. *International Conference on Learning Representations*, 2017.
- Nathan Lambert, Valentina Pyatkin, Jacob Morrison, LJ Miranda, Bill Yuchen Lin, Khyathi Chandu, Nouha Dziri, Sachin Kumar, Tom Zick, Yejin Choi, et al. Rewardbench: Evaluating reward models for language modeling. *arXiv preprint arXiv:2403.13787*, 2024.
- Hunter Lang, David Sontag, and Aravindan Vijayaraghavan. Theoretical analysis of weak-to-strong generalization. *Advances in Neural Information Processing Systems*, 2024.
- Harrison Lee, Samrat Phatale, Hassan Mansoor, Kellie Lu, Thomas Mesnard, Colton Bishop, Victor Carbune, and Abhinav Rastogi. Rlaif: Scaling reinforcement learning from human feedback with ai feedback. *arXiv preprint arXiv:2309.00267*, 2023.
- Harrison Lee et al. Rlaif vs. rlhf: Scaling reinforcement learning from human feedback with ai feedback. In *Forty-first International Conference on Machine Learning*, 2024.
- Kimin Lee, Laura Smith, and Pieter Abbeel. Pebble: Feedback-efficient interactive reinforcement learning via relabeling experience and unsupervised pre-training. In *International Conference on Machine Learning*, 2021.

- Ming Li, Yong Zhang, Shwai He, Zhitao Li, Hongyu Zhao, Jianzong Wang, Ning Cheng, and Tianyi Zhou. Superfiltering: Weak-to-strong data filtering for fast instruction-tuning. *arXiv preprint arXiv:2402.00530*, 2024a.
- Xian Li, Ping Yu, Chunting Zhou, Timo Schick, Luke Zettlemoyer, Omer Levy, Jason Weston, and Mike Lewis. Self-alignment with instruction backtranslation. In *The Twelfth International Conference on Learning Representations*, 2024b.
- Xuechen Li, Tianyi Zhang, Yann Dubois, Rohan Taori, Ishaan Gulrajani, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. AlpacaEval: An automatic evaluator of instruction-following models. https://github.com/tatsu-lab/alpaca_eval, 2023.
- Tianlin Liu, Shangmin Guo, Leonardo Bianco, et al. Decoding-time realignment of language models. In *Proceedings of the 41st International Conference on Machine Learning*, pp. 31015–31031, 2024a.
- Tianqi Liu, Yao Zhao, Rishabh Joshi, Misha Khalman, Mohammad Saleh, Peter J Liu, and Jialu Liu. Statistical rejection sampling improves preference optimization. *International Conference on Learning Representations*, 2024b.
- Yuejiang Liu and Alexandre Alahi. Co-supervised learning: Improving weak-to-strong generalization with hierarchical mixture of experts. *arXiv preprint arXiv:2402.15505*, 2024.
- William Muldrew, Peter Hayes, Mingtian Zhang, and David Barber. Active preference learning for large language models. In *Proceedings of the 41st International Conference on Machine Learning*, pp. 36577–36590, 2024.
- Remi Munos et al. Nash learning from human feedback. In *Forty-first International Conference on Machine Learning*, 2024.
- Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, et al. Webgpt: Browser-assisted question-answering with human feedback. *arXiv preprint arXiv:2112.09332*, 2022.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35: 27730–27744, 2022.
- Arka Pal et al. Smaug: Fixing failure modes of preference optimisation with dpo-positive. *arXiv preprint arXiv:2402.13228*, 2024.
- Gokul Puthumanai, Manav Vora, Pranay Thangeda, and Melkior Ornik. A moral imperative: The need for continual superalignment of large language models. *arXiv preprint arXiv:2403.14683*, 2024.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2023.
- Sayak Ray Chowdhury, Anush Kini, and Nagarajan Natarajan. Provably robust dpo: Aligning language models with noisy feedback. In *Proceedings of the 41st International Conference on Machine Learning*, pp. 42258–42274, 2024.
- Swarnadeep Saha, Omer Levy, Asli Celikyilmaz, Mohit Bansal, Jason Weston, and Xian Li. Branch-solve-merge improves large language model evaluation and generation. *arXiv preprint arXiv:2310.15123*, 2023.
- Charlie Victor Snell, Ilya Kostrikov, Yi Su, Sherry Yang, and Sergey Levine. Offline rl for natural language generation with implicit language q learning. In *The Eleventh International Conference on Learning Representations*, 2023.

- Feifan Song, Bowen Yu, Minghao Li, Haiyang Yu, Fei Huang, Yongbin Li, and Houfeng Wang. Preference ranking optimization for human alignment. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pp. 18990–18998, 2024.
- Nisan Stiennon, Long Ouyang, Jeffrey Wu, Daniel Ziegler, Ryan Lowe, Chelsea Voss, Alec Radford, Dario Amodei, and Paul F Christiano. Learning to summarize with human feedback. *Advances in Neural Information Processing Systems*, 33:3008–3021, 2020.
- Zhiqing Sun, Longhui Yu, Yikang Shen, Weiyang Liu, Yiming Yang, Sean Welleck, and Chuang Gan. Easy-to-hard generalization: Scalable alignment beyond human supervision. *arXiv preprint arXiv:2403.09472*, 2024.
- Yunhao Tang et al. Generalized preference optimization: A unified approach to offline alignment. *ICML*, 2024.
- Antti Tarvainen and Harri Valpola. Mean teachers are better role models: Weight-averaged consistency targets improve semi-supervised deep learning results. *Advances in Neural Information Processing Systems*, 2017.
- Gemini Team, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, et al. Gemini: a family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805*, 2023.
- Gemma Team, Thomas Mesnard, Cassidy Hardin, Robert Dadashi, Surya Bhupatiraju, Shreya Pathak, Laurent Sifre, Morgane Rivière, Mihir Sanjay Kale, Juliette Love, et al. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295*, 2024.
- Ankur P Parikh Thibault Sellam, Dipanjan Das. Bleurt: Learning robust metrics for text generation. In *Proceedings of ACL*, 2020.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shrutu Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- Qizhe Xie, Minh-Thang Luong, Eduard Hovy, and Quoc V Le. Self-training with noisy student improves imagenet classification. In *Computer Vision and Pattern Recognition*, pp. 10687–10698, 2020.
- Wei Xiong, Hanze Dong, Chenlu Ye, Han Zhong, Nan Jiang, and Tong Zhang. Gibbs sampling from human feedback: A provable kl-constrained framework for rlhf. *arXiv preprint arXiv:2312.11456*, 2023.
- Rui Yang et al. Rewards-in-context: Multi-objective alignment of foundation models with dynamic preference adjustment. In *Forty-first International Conference on Machine Learning*, 2024.
- Seonghyeon Ye, Doyoung Kim, Sungdong Kim, Hyeonbin Hwang, Seungone Kim, Yongrae Jo, James Thorne, Juho Kim, and Minjoon Seo. FLASK: Fine-grained language model evaluation based on alignment skill sets. In *International Conference on Learning Representations*, 2024.
- Min-Hsuan Yeh, Leitian Tao, Jeffrey Wang, Xuefeng Du, and Yixuan Li. How reliable is human feedback for aligning large language models? *arXiv preprint arXiv:2410.01957*, 2024.
- Meng Yu et al. Simpo: Simple preference optimization with a reference-free reward. *arXiv preprint arXiv:2405.14734*, 2024.
- Weizhe Yuan, Richard Yuanzhe Pang, Kyunghyun Cho, Sainbayar Sukhbaatar, Jing Xu, and Jason Weston. Self-rewarding language models. *arXiv preprint arXiv:2401.10020*, 2024.
- Zheng Yuan, Hongyi Yuan, Chuanqi Tan, Wei Wang, Songfang Huang, and Fei Huang. Rrhf: Rank responses to align language models with human feedback without tears. *arXiv preprint arXiv:2304.05302*, 2023.
- Yongcheng Zeng, Guoqing Liu, Weiyu Ma, et al. Token-level direct preference optimization. In *Proceedings of the 41st International Conference on Machine Learning*, pp. 58348–58365, 2024.

Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, Christopher Dewan, Mona Diab, Xian Li, Xi Victoria Lin, et al. Opt: Open pre-trained transformer language models. *arXiv preprint arXiv:2205.01068*, 2022.

Yao Zhao et al. Slic-hf: Sequence likelihood calibration with human feedback. *arXiv preprint arXiv:2305.10425*, 2023.

Chujie Zheng, Ziqi Wang, Heng Ji, Minlie Huang, and Nanyun Peng. Weak-to-strong extrapolation expedites alignment. *arXiv preprint arXiv:2404.16792*, 2024.

Zhanhui Zhou, Zhixuan Liu, Jie Liu, Zhichen Dong, Chao Yang, and Yu Qiao. Weak-to-strong search: Align large language models via searching over small language models. *Advances in Neural Information Processing Systems*, 2024.

Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.

A EXPERIMENTAL DETAIL

A.1 TRAINING DETAILS

Software and hardware. Our experiments are conducted on servers equipped with NVIDIA A100 GPUs, with 80 GB of VRAM. The operating system used is Ubuntu 22.04.2 LTS, supported by NVIDIA CUDA Toolkit version 12.1 and cuDNN version 8.9. All experimental implementations are carried out in Python version 3.11.4, utilizing the PyTorch framework version 1.12.1.

Training hyper-parameters. Based on *TRL*², we deploy the training of both teacher and student models with the same hyperparameters as shown in Table 7 and Table 8.

Reddit TL;DR dataset preprocessing details. For the Reddit TL;DR dataset, we preprocess the data by filtering out samples with fewer than 512 tokens. This results in 92,000 training samples and 8,000 test samples. We divide the training data into two disjoint sets of 46,000 samples each. The first set serves as the labeled dataset, denoted as $\mathcal{D}_l$, while the second set, referred to as the unlabeled dataset, $\mathcal{D}_u$, disregards the preference labels.

A.2 EVALUATION DETAILS

Table 6: Sample prompt for the GPT-4 evaluation. Text highlighted in **orange** represents the question, while text in **blue** represents the responses under comparison.

System prompt [System] You are a helpful, harmless, and precise assistant for checking the quality of the answer. We would like to request your feedback on the performance of two AI assistants in response to the user question displayed above. Please rate the harmlessness, helpfulness, and level of detail of their responses. Your evaluation should consider factors such as the helpfulness, harmlessness, relevance, accuracy, depth, creativity, and level of detail of the response. Each assistant receives an overall score on a scale of 1 to 10, where a higher score indicates better overall performance. Please first output a single line containing only two values indicating the scores for Assistant 1 and 2, respectively. The two scores are separated by a space. In the subsequent line, please provide a comprehensive explanation of your evaluation, avoiding any potential bias and ensuring that the order in which the responses were presented does not affect your judgment. User prompt [Question] ###Human: What are some symptoms of caffeine withdrawal? ###Assistant: [The Start of Assistant 1's Answer] When caffeine is removed from someone who has habitually consumed it, the body has a natural reaction to produce more adrenaline, in order to reduce the need for caffeine. This increased adrenaline production causes many different symptoms, ranging from jitteriness and restlessness, to lethargy and drowsiness. [The End of Assistant 1's Answer] [The Start of Assistant 2's Answer] You might feel an overall sense of mental tiredness, a decreased ability to concentrate, and problems sleeping. You may also experience headaches and a sense of jitteriness or edginess. There are also some physical symptoms that can appear, such as muscle pain and vomiting. [The End of Assistant 2's Answer]

GPT-4 evaluation details. Table 6 presents the prompts and responses in our GPT-4 evaluation, adopted from (Khanov et al., 2024). Each GPT-4 request comprises both a system and a user prompt. The system prompt delineates the proxy’s attributes and its specific task, while the user prompt poses a question and provides responses from the two methods.

Gold reward model. For the evaluation of the HH-RLHF, we leverage the state-of-the-art gold reward model Ray2333/reward-model-Mistral-7B-instruct-Unified-Feedback from the RewardBench leaderboard (Lambert et al., 2024), which fine-tunes Mistral-7B-Instruct-v0.2 on the llm-blender/Unified-Feedback dataset. We also verify our results consistently hold

²<https://github.com/huggingface/trl>

under different choices of gold reward models in Appendix C. For Reddit TL;DR, we take the state-of-the-art gold reward model `OpenAssistant/reward-model-deberta-v3-large-v2` on this dataset from the RewardBench leaderboard (Lambert et al., 2024) for evaluation.

Hyper-parameters for model generation. To evaluate the response of the model, we adopted the temperature as 0.7 and the max tokens as 256.

Table 7: Summary of training hyperparameters for supervised fine-tuning and direct preference optimization for OPT-family models for HH-RLHF.

	Parameters	Value
Supervised fine-tuning	Number of epochs	1
	Learning rate	1×10^{-5}
	Batch size	32
	Gradient accumulation steps	1
	Maximum sequence length	512
	DeepSpeed Zero stage	2
	LoRA rank	0
Direct preference optimization	Number of epochs	1
	Learning rate	5×10^{-5}
	β	0.1
	Batch size	16
	Gradient accumulation steps	1
	Maximum sequence length	512
	DeepSpeed Zero stage	2
	LoRA rank	8

Table 8: Summary of training hyperparameters for supervised fine-tuning and direct preference optimization for OPT-family models on Reddit TL;DR.

	Parameters	Value
Supervised fine-tuning	Number of epochs	1
	Learning rate	1×10^{-5}
	Batch size	32
	Gradient accumulation steps	1
	Maximum sequence length	512
	DeepSpeed Zero stage	2
	LoRA rank	0
Direct preference optimization	Number of epochs	1
	Learning rate	5×10^{-6}
	β	0.5
	Batch size	16
	Gradient accumulation steps	1
	Maximum sequence length	512
	DeepSpeed Zero stage	2
	LoRA rank	8

B ADDITIONAL QUALITATIVE EXAMPLES

In Table 9, we provide additional qualitative examples from $\mathcal{D}_{\text{weak}}$ where the chosen response by weak LLM mismatches that of the human annotators. Despite contradicting the human preference, the gold reward of the response chosen by the weak LLM can be higher, i.e., $r_{\text{gold}}(x, \hat{y}_c) > r_{\text{gold}}(x, \hat{y}_r)$. This indicates that the preferences of the weak LLM can align with the gold reward model, even if they differ from human labels.

Table 9: Comparative examples of the chosen and rejected responses in $\mathcal{D}_{\text{weak}}^{\text{mismatch}}$.

Example 1 Prompt Human: How does a pacemaker help the function of the heart? Assistant: Chosen response by weak LLM (rejected by human annotators) A pacemaker is a medical device implanted in the chest which stimulates the heart to beat at a normal rhythm, helping the heart to pump and circulate blood more efficiently and effectively. Pacemakers are generally implanted with catheters by a practitioner in a cardiac catheterization or angiography procedure, or in rare cases with open chest surgery. Pacemaker systems usually consist of a pulse generator and set of leads, one end of which is sutured to the heart (lead to the right atrium, right ventricle, right-side of the chest, or the right coronary sinus), and the other end which is connected to the pacemaker. Pacemakers also can be used as a lifesaving measure for patients with heart failure, or at risk of permanent damage from their heart not functioning properly. (gold reward: 11.99) Rejected response by weak LLM (chosen by human annotators) A pacemaker is a small electronic device implanted under the skin with an electrode attached to the heart to help control the rate and rhythm of the heartbeat. The pacemaker senses when the heartbeat is slow or irregular and sends electrical impulses to the heart at a pre-defined safe and healthy rate, normalizing the heart beat. (gold reward: 11.76)
Example 2 Prompt Human: Why does someones weight affect their tolerance to alcohol? Assistant: Chosen response by weak LLM (rejected by human annotators) Studies suggest that people with a higher body weight may have a lower tolerance for alcohol, but this is not well understood. It may be due to differences in body composition or other factors that influence how alcohol is metabolized and how it affects the central nervous system. (gold reward: 8.05) Rejected response by weak LLM (chosen by human annotators) Well, alcohol is stored as fat in the human body, and so one's tolerance to alcohol is related to one's fat stores. (gold reward: 6.66)
Example 3 Prompt Human: How does a smart grid work for powering a country's electric grid? Assistant: Chosen response by weak LLM (rejected by human annotators) So we're talking about electricity grids, right? One way to describe what's happening is to say that the grid acts as a network for power distribution, where energy flows from various energy sources to various consumers. And it's a bit like a huge complex network of pipes and nodes. But in this case, the 'nodes' are basically power plants that create and distribute electricity, and the 'pipes' are the electrical lines which transport power from the sources to the nodes. (gold reward: 2.38) Rejected response by weak LLM (chosen by human annotators) There's a bunch of different ways to generate electricity. Most of the time, they boil water to produce steam, which spins a turbine, which spins a generator, which generates electricity. But there are other ways to do it. To start, I need to understand what you mean by 'smart grid'. Are you talking about a grid that's connected to a smart city, or a smart grid in terms of power lines, or something else? (gold reward: 1.07)

C ADDITIONAL EXPERIMENTS

Evaluation with more gold reward models. To verify the reliability of our conclusions, we employ different gold reward models to evaluate the quality of model output. Specifically, we leverage another two competitive gold reward models weqweasdas/RM-Mistral-7B and OpenAssistant/reward-model-deberta-v3-large-v2 from the RewardBench leaderboard (Lambert et al., 2024). As shown in Figure 5, using these two alternative gold reward models,

we can still observe that alignment with weak LLM feedback can outperform human feedback. This confirms the reliability of our findings, and its insensitivity to the choice of gold reward model.




Figure 5: (a) Evaluation with the gold reward model reward-model-deberta-v3-large-v2. (b) Evaluation with the gold reward model RM-Mistral-7B.

Results on additional weak LLM feedback. To validate our conclusion that weak LLM provides feedback comparable to human feedback, we extend our analysis beyond the specific case of OPT-125M. We incorporated GPT-Neo-125M (Gao et al., 2020) and Pythia-160M (Biderman et al., 2023) as additional weak LLM to provide feedback. We use OPT-1.3B and Mistral-7B as student models to be aligned. Following the alignment with weak LLM feedback process outlined in Section 3.1, we assess the generalizability of our conclusion. As shown in Table 10, models aligned with feedback from weak LLMs GPT-Neo-125M and Pythia-160M achieved performance comparable to the model aligned with human feedback. This demonstrates that our conclusion is not dependent on a single weak model, but rather applies more broadly across different weak LLMs.

Table 10: Results on using different weak LLM supervisors: OPT-125M, GPT-Neo-125M and Pythia-160M.

	Human feedback	OPT-125M	GPT-Neo-125M	Pythia-160M
OPT-1.3B	4.63	4.84	4.55	4.87
Mistral-7B	7.19	7.90	7.27	8.01

Correlation of generation quality. In Figure 6, we investigate whether the generation quality correlates between models using weak LLM feedback vs. human feedback. In particular, for each test input prompt, we measure the sentence-level similarity between the model’s generation $\hat{y}$ (using either π_θ^* or π_h^*) and the chosen response y_w by human annotators from HH-RLHF. Thibault Sellam (2020) proposed a cosine similarity based on the BLUERT embedding. These similarities are denoted as $\text{sim}_{\pi_\theta^*}(\hat{y}, y_c)$ for the model aligned with weak LLM feedback and $\text{sim}_{\pi_h^*}(\hat{y}, y_c)$ for the model aligned with human feedback, represented on the x -axis and y -axis respectively. We observe a moderate correlation between the two across all test samples from HH-RLHF.

Additionally, we examine the correlation between the model π_θ^* aligned with weak LLM feedback and the weak LLM itself π_w^* on the right of Figure 6. This analysis helps reveal whether aligning a model under weak feedback could risk imitating the errors inherent in a weak LLM. We observe that the correlation between π_θ^* and π_w^* (with $R^2 = 0.4888$) is relatively weaker compared to that between π_θ^* and π_h^* (with $R^2 = 0.5789$). These results suggest that the model π_θ^* not only aligns more closely with π_h^* , but also effectively extrapolates beyond imitating the weaker teacher.

Ablation on the impact of SFT. Learning from preference data typically begins by fine-tuning a pre-trained language model with supervised learning on high-quality data for the downstream tasks of interest. We further explore the impact of SFT when aligning the target policy model (*c.f.* Equation 6). We utilize the OPT-125M as the weak LLM to provide feedback and the OPT-1.3B as the student model, conducting our experiments on the HH-RLHF dataset. For the model π_θ^* aligned with weak LLM feedback and the model π_h^* aligned with human feedback, we compare the following two settings: (1) **DPO (w/o SFT)**: Directly using the pre-trained model as the reference model or




Figure 6: **Left:** Agreement between models using weak LLM feedback (π_θ^*) vs. human feedback (π_h^*). **Right:** Agreement between the model (π_θ^*) aligned with weak LLM feedback and the weak LLM itself (π_w^*).

initialization. (2) **DPO (w SFT)**: Use fine-tuned model π_θ^{SFT} as a reference model for DPO training. Compared to directly using the untuned base model as a reference model, performing SFT enhances the model’s ability to generate desired responses to questions.

Table 11: Ablation on the impact of supervised fine-tuning (SFT) as initialization for DPO. We report the average gold reward. Numbers in brackets denote improvement relative to the pre-trained model (OPT-1.3B) without any fine-tuning.

Method	Weak LLM feedback	Human feedback
DPO (w/o SFT)	3.83 (+1.61)	3.77 (+1.55)
DPO (w SFT)	4.84 (+2.61)	4.63 (+2.41)

Ablation on the KL coefficient (β) for weak LLM feedback for alignment. The β parameter in Equation 3 functions as the KL coefficient during DPO training, with higher values indicating more stringent regularization. We employ OPT-125M as the weak LLM to provide feedback and OPT-1.3B as the student model, performing alignment using varying β values: $\{0.05, 0.1, 0.2, 0.3, 0.5\}$ on the HH-RLHF dataset. As shown in the left of Figure 7, the gold reward for the model alignment with weak LLM feedback (π_θ^*) is high under relatively small β values such as 0.05 and 0.1, and starts to decline as β increases. A similar trend is observed for the model trained with human feedback.




Figure 7: Ablation on the KL Coefficient (β) for alignment with weak LLM feedback and human feedback.

Dataset division with different random seeds. We randomly partition the entire training set into a labeled dataset, denoted as $\mathcal{D}_l$, and an unlabeled dataset, $\mathcal{D}_u$, using a random seed of 22. To ensure that our conclusions are insensitive to the data division, we also utilized five additional random seeds

to perform the experiments. Same as our main experiment, we employ OPT-125M as the weak LLM and OPT-1.3B as the student model on the HH-RLHF dataset. Across 5 random runs, the average gold rewards for the π_θ^* and π_h^* are 4.81 and 4.67, respectively, with variances of 0.12 and 0.14. These results affirm that our conclusions are statistically stable under different dataset splits.

Understanding the different impact between weak LLM feedback vs. random noise. Our investigation in Section 4 suggests that weak LLM feedback contains more systematic noise (e.g., errors primarily occur between choices that are subtly different). In this ablation, we contrast performing alignment with feedback from weak LLM by using the feedback containing *random noise*. Adopting OPT-125M as the weak LLM and OPT-1.3B as the student model, we design three controlled settings for this ablation:

- $\mathcal{D}_{\text{random}}^{\text{match}}$: Randomly sample 60.6% of the samples from $\mathcal{D}_u$, and assign human feedback, i.e., $(r_w(x, y_c) > r_w(x, y_r))$. This set has the same size as our weakly supervised set $\mathcal{D}_{\text{weak}}^{\text{match}}$.
- $\mathcal{D}_{\text{random}}^{\text{mismatch}}$: Use the remainder 39.4% samples from $\mathcal{D}_u$, and assign preference labels *opposite to the human feedback*, i.e., $(r_w(x, y_c) < r_w(x, y_r))$.
- $\mathcal{D}_{\text{random}}$: The union of $\mathcal{D}_{\text{random}}^{\text{match}}$ and $\mathcal{D}_{\text{random}}^{\text{mismatch}}$.

As demonstrated in Table 12, the student model trained with the match set $\mathcal{D}_{\text{random}}^{\text{match}}$ outperforms the model trained with $\mathcal{D}_{\text{random}}$ by 0.66 in terms of gold reward. This suggests that removing random noise in preference datasets can markedly enhance alignment performance. Conversely, the performance of the model trained with $\mathcal{D}_{\text{random}}^{\text{mismatch}}$ exhibits a gold reward that is significantly lower (1.53), illustrating the negative effect of random noise on model alignment. These findings highlight the difference between random noise and the “noise” from weak LLM feedback.

Table 12: Effect of purifying random noise. Numbers in brackets denote improvement relative to the pre-trained student model without any fine-tuning.

	% matching human feedback	Gold reward of π_θ^*
$\mathcal{D}_{\text{random}}$	60.6%	3.91 (+1.69)
$\mathcal{D}_{\text{random}}^{\text{match}}$	100%	4.57 (+2.35)
$\mathcal{D}_{\text{random}}^{\text{mismatch}}$	0%	1.53 (-0.69)

Exploring the impact of DPO-trained weak LLMs vs. traditional reward models. To maintain a consistent and unified framework for training both the teacher and student models, we adopt DPO instead of traditional reward modeling for teacher-model training. Furthermore, since DPO is inherently formulated as a Bradley-Terry loss, it naturally functions as an implicit reward model. When the weak LLM’s role is limited to generating reward signals, it is also feasible to replace it with a traditional Bradley-Terry reward model. To explore this alternative, we conducted additional experiments comparing DPO-trained weak LLMs with traditional reward models. The results show that both approaches achieve comparable performance, with both surpassing human labelers.

Table 13: Comparison of gold reward between DPO-trained weak LLMs and traditional reward models.

Feedback Source	Gold Reward
Weak LLM (with DPO)	4.84
Weak LLM (with reward modeling)	4.86
Human labeler	4.63

D LIMITATIONS AND FUTURE WORK

Limitations. Our research presents an in-depth investigation of alignment with the feedback from weak LLM. We delve deeply into the reasons underlying this phenomenon. However, our study is not without limitations, which we aim to address in subsequent research. Our focus has predominantly been on empirical experiments, leaving the theoretical underpinnings of our findings less explored. These areas present opportunities for further exploration and development.

Implications for future alignment research. As we look toward the future of AI alignment, several recommendations emerge from our study that could further refine the practice and enhance the reliability of alignment methodologies. First, integrating hybrid feedback systems that combine human insights with AI-generated feedback could leverage the strengths of both, minimizing the limitations inherent in each approach. Secondly, it is crucial to develop more sophisticated metrics for evaluating alignment quality that go beyond traditional accuracy metrics, to capture the nuanced understanding and generative capabilities required in real-world applications. Finally, exploring the ethical implications of AI-generated feedback and ensuring that these systems adhere to ethical guidelines is vital. By addressing these areas, the field can move towards more effective, efficient, and ethically responsible AI alignment strategies that are capable of supporting the safe integration of AI systems into society.